

TP3_apartado_2_kafka_cluster-adrian_zapater

January 11, 2020

1 TP3 - Adrián José Zapater Reig

1.1 Apartado 2 - Kafka + Spark streaming

Primero levantamos el cluster de kafka:

```
[ ]: # Zookeeper en el puerto 2181
! /usr/local/kafka/bin/zookeeper-server-start.sh /usr/local/kafka/config/
↪zookeeper.properties
```

```
[2]: # Broker1 puerto 9092
! cat /usr/local/kafka/config/broker_3.properties
```

```
broker.id=0
listeners=PLAINTEXT://:9092
log.dirs=/tmp/kafka-logs-0
```

```
num.network.threads=3
num.io.threads=8
```

```
socket.send.buffer.bytes=102400
socket.receive.buffer.bytes=102400
socket.request.max.bytes=104857600
```

```
num.partitions=1
num.recovery.threads.per.data.dir=1
```

```
offsets.topic.replication.factor=1
transaction.state.log.replication.factor=1
transaction.state.log.min.isr=1
```

```
log.retention.hours=168
log.segment.bytes=1073741824
log.retention.check.interval.ms=300000
```

```
zookeeper.connect=localhost:2181
zookeeper.connection.timeout.ms=6000
```

```
group.initial.rebalance.delay.ms=0
```

```
[ ]: ! /usr/local/kafka/bin/kafka-server-start.sh /usr/local/kafka/config/broker_3.  
      ↪ properties
```

```
[3]: # Broker2 puerto 9093  
      ! cat /usr/local/kafka/config/broker_1.properties
```

```
broker.id=1  
listeners=PLAINTEXT://:9093  
log.dirs=/tmp/kafka-logs-1
```

```
num.network.threads=3  
num.io.threads=8
```

```
socket.send.buffer.bytes=102400  
socket.receive.buffer.bytes=102400  
socket.request.max.bytes=104857600
```

```
num.partitions=1  
num.recovery.threads.per.data.dir=1
```

```
offsets.topic.replication.factor=1  
transaction.state.log.replication.factor=1  
transaction.state.log.min.isr=1
```

```
log.retention.hours=168  
log.segment.bytes=1073741824  
log.retention.check.interval.ms=300000
```

```
zookeeper.connect=localhost:2181  
zookeeper.connection.timeout.ms=6000
```

```
group.initial.rebalance.delay.ms=0
```

```
[ ]: ! /usr/local/kafka/bin/kafka-server-start.sh /usr/local/kafka/config/broker_1.  
      ↪ properties
```

```
[4]: # kafkaTwitter topic usando broker o como bootstrap server.  
      ! /usr/local/kafka/bin/kafka-topics.sh --create --bootstrap-server localhost:  
      ↪ 9092 --replication-factor 2 --partitions 2 --topic kafkaTwitter
```

```
[5]: 1
```

```
__consumer_offsets
```

kafkaTwitter